

An Open Firm-Level Microsimulation for UK VAT Threshold Policy Analysis*

Vahid Ahmadi[†]

July 2026

Abstract

I present an open firm-level microsimulation for conditional costing of UK VAT registration-threshold reforms. A weighted synthetic population is calibrated to ONS and HMRC aggregates on two declared universes, and level, taper and reduced-rate scenarios are costed on a common base. With membership evaluated on data-year turnover, a mechanical costing of the April 2024 rise from £85,000 to £90,000 lies within a third of HMRC's published profile in each release year with no tuned parameter. The notch's dominated region is £21,250 for a firm with no reclaimable inputs but, once unregistered firms bear input VAT, firm-specific and zero for about half of near-threshold firms; a reduced-rate band then creates a distortion where none existed. Intensive-margin responses are second order. The near-threshold density is target-inherited, so the data cannot support bunching inference.

Keywords: value-added tax, microsimulation, tax notch, registration threshold

JEL codes: H25, H26, C63

***Acknowledgements:** I thank Max Ghenis and Seyed Mahdi Hosseini for their great suggestions and helpful discussions.

[†]Research Associate at PolicyEngine. Email: vahid@policyengine.org.

1 Introduction

The United Kingdom’s Value Added Tax (VAT) applies above a registration threshold that is the highest in the OECD alongside Switzerland (Seely, 2024; HM Revenue and Customs, 2024). The threshold is a *notch* rather than a kink: once a firm’s annual taxable turnover crosses it, VAT falls due on the firm’s whole base, not only on the part above the threshold, so registration produces a discrete jump in liability rather than a marginal one. The threshold falls in a dense part of the firm distribution. Figure 1(a) shows that VAT- and PAYE-registered UK businesses are concentrated at the lower end of the turnover distribution, and Figure 1(b) shows that the VAT-registered population per £1,000 of turnover is densest at and below the threshold, so the notch applies to a large number of firms.




Figure 1: The UK firm-size distribution and VAT registration, 2023–24 (threshold £85,000). Panel (a): number of businesses by turnover band in the ONS business register, which covers businesses registered for VAT and/or PAYE (about 2.7 million); the full private-sector business population including the 2.9 million businesses registered for neither is about 5.5 million (DBT business population estimates, start of 2024). Panel (b): number of VAT-registered firms by turnover band. Band widths are unequal, so cross-band count comparisons should be read per £1,000 of turnover. Source: panel (a), ONS, *UK Business: Activity, Size and Location*, 2024 edition (Table 8, enterprises by turnover sizeband); panel (b), HMRC, *Value Added Tax (VAT) Annual Statistics* (VAT-registered population by turnover band).

Firms respond to the notch by locating below the threshold. Tabulations of HMRC turnover data show the number of firms rising towards £84,000 and then dropping sharply at the £85,000 threshold, leaving excess mass just below it (Office for Budget Responsibility, 2023). Liu et al. (2021) document this bunching in UK administrative firm data over 2004–05 to 2014–15, when the threshold rose from £58,000 to £81,000, and Liu et al. (2024) find that firms slow their turnover growth as they approach the threshold. The

notch therefore distorts the size distribution and the growth decisions of firms near the margin, which places the threshold under recurrent pressure for reform.

Reform proposals change the threshold in one of three ways: its *level* (raising or lowering it), its *shape* (replacing the notch with a schedule that phases the rate in over a band), or its *rate* (a reduced rate for firms in a band above the threshold). The three are not equivalent. Changing the level moves the notch along the turnover axis; changing the shape or the rate alters the size of the discrete jump itself. Costing the three on a common basis is the practical difficulty this paper addresses. A change in the level can be costed by mechanical accounting, but a change in the shape of the schedule cannot be evaluated with the reduced-form bunching methods used to study the threshold, because a continuous schedule leaves no single point of excess mass to exploit. The firm-level data needed to study turnover near the threshold is, in addition, not available in open form.

In this paper I build an open firm-level accounting microsimulation for conditional costing of UK VAT threshold reforms on a common static basis. I construct weighted synthetic firm records calibrated to published official aggregates—HMRC registered-firm counts by turnover band and by sector, net VAT liability by turnover band, and ONS business counts and employment bands—for the 2023–24 tax year, when the threshold was £85,000, and separately for 2024–25, when it rose to £90,000. These are not reconstructed administrative records: turnover and liability within published bands are simulated under explicit assumptions. The released code permits those assumptions and counterfactual schedules to be changed and audited.

On these data I do four things. First, I cost level reforms and schedule reforms—a graduated taper and a reduced-rate band—statically, by applying each counterfactual schedule mechanically to the firm distribution. Second, I characterise the notch’s distortion through its *dominated region*—the exact, elasticity-free range of turnover above the threshold in which no firm profits from locating—and how each reform changes it. Third, I price the level and rate reforms behaviourally: each firm re-optimises its turnover within its schedule region, conditional on an assumed turnover elasticity e that the data do not identify, reported as an e -sensitivity range that nests the static costing exactly in its $e \rightarrow 0$ limit. Fourth, I show that the synthetic data’s near-threshold structure is target-inherited and cannot support bunching inference: the £85,000 profile is calibrated to the OBR’s published £1,000-band data and the generator contains no location-choice mechanism.

Four findings follow. First, the static pipeline prices level, shape, and rate reforms on one basis; its anchor—the April 2024 rise from £85,000 to £90,000, whose mechanical cost profile lands within a third of HMRC’s published costing in each of the four release years (within 10% in 2025–26 and 2026–27) and 36% above it in the sign-flip year, with no retention or behavioural parameter tuned—is a consistency check rather than an out-of-sample validation, since the population is calibrated to HMRC aggregates (Section 4). Second, the shape and rate reforms act on the notch’s distortion in a way a level change does not: the notch creates an exact £21,250 dominated region for a firm bearing no unreclaimed input VAT—the upper bound; once unregistered firms bear VAT on inputs the region is firm-specific and vanishes for the 46% of near-threshold firms whose standard-rated input share exceeds one half—which a rise relocates and widens, a reduced-rate band splits

across two notches, and only a continuous taper removes—and the taper’s band width is a property of its shape, not a requirement of removal (Section 6). Third, conditional on the assumed elasticity, intensive-margin responses barely move the costings (Section 7). Fourth, the synthetic data carry a transparency lesson: their near-threshold excess mass is inherited from calibration targets or band edges, is unstable across specifications, and cannot validate a bunching estimator or identify behaviour (Section 5); the behavioural evidence remains the administrative-data work of [Liu et al. \(2021\)](#). Table 4 collects the headline estimates.

Related literature

The paper draws on four strands of work. The first is the methodology of bunching and notches ([Chetty et al., 2011](#); [Saez, 2010](#); [Kleven and Waseem, 2013](#); [Kleven, 2016](#); [Best et al., 2015](#)), which provides the analytic notion of a dominated region; I use the dominated region to characterise a distortion that leaves no single point of excess mass to exploit. Recent evidence from a Finnish payroll-tax notch shows that a policy threshold can impede mobility and firm growth well beyond the immediate cutoff and can bias conventional difference-in-differences comparisons ([Jysmä et al., 2026](#)). That result strengthens the case for treating my local dominated-region metric as partial: its payroll and employment setting does not identify the response to the UK VAT turnover threshold, but it demonstrates the potentially wider dynamic margin that this accounting model omits.

The second is the empirical literature on VAT registration thresholds ([Liu et al., 2021, 2024](#); [Onji, 2009](#); [Harju et al., 2019](#); [Asatryan and Peichl, 2017](#); [Nandi and Warwick, 2020](#); [Bashir et al., 2024](#)). [Liu et al. \(2021\)](#) is closest: it estimates bunching at the UK notch on administrative data over 2004–2014 and documents voluntary registration; I take its estimates as the established behavioural facts and cost shape and rate reforms statically on a common basis, which that literature does not do. The international evidence ([Harju et al., 2019](#); [Onji, 2009](#); [Asatryan and Peichl, 2017](#); [Nandi and Warwick, 2020](#)) notes that bunching at a VAT threshold can reflect compliance costs, firm-splitting, or reporting rather than real turnover responses.

The third strand questions whether bunching identifies a structural elasticity. [Blomquist et al. \(2021\)](#) show that it is not non-parametrically identified at a kink or notch without strong functional-form assumptions; [Bertanha et al. \(2023\)](#) develop the notch case. I therefore do not infer an elasticity from the synthetic data. The fourth strand is the theory of size-based tax thresholds: the optimal VAT threshold of [Keen and Mintz \(2004\)](#), thresholds as partitions of compliance in [Kanbur and Keen \(2014\)](#), and optimal remittance with firm-level administrative costs in [Dharmapala et al. \(2011\)](#), together with the structural literature on size-dependent policy ([Garicano et al., 2016](#); [Gourio and Roys, 2014](#)); the Mirrlees Review’s VAT chapter frames the UK debate ([Crawford et al., 2010](#)). The present paper makes no comparable structural claim. Firm-level tax microsimulation has precedents in corporate-tax models such as the German BizTax model ([Fossen and Bach, 2008](#)); to my knowledge no open firm-level model of the UK VAT threshold exists, which is the gap this paper fills.

This paper makes three contributions. First, it provides an open and auditable firm-

level accounting microsimulation for conditional VAT-threshold scenarios, where the existing behavioural evidence rests on restricted administrative data. Second, it characterises the notch through its model-implied dominated region and costs two schedule reforms—a graduated taper and a reduced-rate band—on the same static basis as level reforms. Third, it demonstrates that band-calibrated synthetic data carry no sub-band information beyond their targets and distributional assumptions. All fiscal estimates below should therefore be read as conditional scenario outputs, not point forecasts with administrative-data precision. Identifying the turnover elasticity from administrative microdata, welfare accounting, revenue-neutral designs, and an explicit model of voluntary registration are left for future work.

The remainder of the paper proceeds as follows. Section 2 describes the institutional background of the UK VAT threshold. Section 3 sets out the synthetic firm microdata and its calibration to official aggregates. Section 4 sets out the static costing of threshold changes and the schedule reforms. Section 5 shows that the synthetic data contain no usable behavioural bunching signal. Section 6 develops the notch and its exact dominated region, and shows how each reform changes it. Section 7 prices the reforms behaviourally, conditional on an assumed elasticity, and derives the level-rise invariance. Section 8 concludes.

2 Institutional background

Registration and the notch. The United Kingdom levies a value-added tax at a standard rate of 20% on most goods and services. A business must register once its taxable turnover over any rolling twelve-month period exceeds the registration threshold, after which it charges VAT on its taxable sales and reclaims VAT paid on inputs. Because registration makes VAT due on the firm’s *entire* taxable turnover rather than only the amount above the cutoff, the threshold is a *notch*—a discrete jump in liability—and not a kink. The cost of crossing is largest for consumer-facing firms with limited pass-through and small reclaimable input shares, for whom the incentive is to restrain turnover and locate just below the threshold (Liu et al., 2021; Kleven and Waseem, 2013).

Value added, input reclaim, and voluntary registration. Real VAT is a tax on *value added*, not on turnover: a registered firm charges output VAT on its sales but reclaims input VAT on its purchases. The model reflects this: each synthetic firm’s net liability is the standard rate applied to its value added (Section 3), and under the value-added formulation—in which unregistered firms bear no unreclaimed input VAT—the notch’s dominated region is invariant to the firm’s input share (Section 6). Two features of the real system remain outside the model’s scope. A large share of below-threshold firms register *voluntarily*—roughly 43% of eligible below-threshold firms (Liu et al., 2021), driven by input-VAT reclaim and business-to-business sales—and for them the threshold carries no bunching incentive; voluntary registration enters the model only as a registration-rate parameter and a retention sensitivity, not as a choice. And net-repayment traders (zero-rated exporters and similar, HMRC’s negative-liability column) are not represented, because the public aggregates do not resolve rated structure by

turnover band.

Threshold history. The registration threshold was held at £85,000 from 1 April 2017 to 31 March 2024—a seven-year nominal freeze, the longest in the history of the tax—before being raised to £90,000 on 1 April 2024, with the deregistration threshold rising from £83,000 to £88,000 (HM Revenue and Customs, 2024; Seely, 2024). Because nominal turnover grows with inflation and real activity, freezing the threshold lowers it in real terms and draws an increasing number of firms towards the registration margin; this fiscal-drag mechanism was a central motivation for the April 2024 increase. My data year is 2023–24, so the prevailing statutory threshold throughout my analysis is £85,000. Figure 2 shows the pattern in the turnover distribution: the number of firms rises towards £84,000 and falls sharply above £85,000, and the projected 2025–26 distribution is more sharply bunched than the 2019–20 outturn as fiscal drag draws more firms towards the frozen threshold.




Figure 2: Bunching below the £85,000 VAT registration threshold: number of businesses by £1,000 turnover band. Source: OBR, *Economic and Fiscal Outlook March 2023*, Chart C (Open Government Licence).

The reform debate. The UK threshold is the highest in the OECD alongside Switzerland (HM Revenue and Customs, 2024), and there is debate over whether to raise it—to reduce the compliance burden on small firms—or lower it, to broaden the base and reduce the growth distortion at the margin. Three families of reform feature, and I evaluate all three: a change in the threshold’s *level* (the £85,000 to £90,000 increase of April 2024 and further moves around it); a change in its *shape* (a *graduated taper* that phases the effective rate in over a band, replacing the notch with a continuous profile); and a change in its *rate* (a *reduced rate* for small firms in a band above the threshold, with the standard 20% rate retained above). The taper and reduced-rate designs follow proposals in the UK policy debate: Office of Tax Simplification (2017) canvassed a financial incentive or reduced effective rate to smooth the cliff-edge at the threshold, and HM Treasury’s

call for evidence on the registration threshold examined smoothing mechanisms in the same family (HM Treasury, 2018). The European Union ran a closely related “graduated relief” scheme for small enterprises and abolished it from January 2025 on complexity grounds, in favour of a simple exemption with a transition tolerance (Council of the European Union, 2020)—relevant precedent for the administrability of taper-style designs. A banded reduced rate introduces a smaller *secondary* notch at the band top, which the continuous taper avoids. I cost all three by the same machinery in Section 4.

3 Data

Studying firm responses at the threshold needs firm-level turnover near the cutoff, which the official statistics publish only in coarse turnover bands. I therefore build a synthetic firm population calibrated to those aggregates: it reproduces selected official totals while supplying explicitly simulated within-band records for accounting exercises. These records are not estimates of the confidential firms that the authorities do not release. The same calibrate-to-published-aggregates approach underlies open household-level tax-benefit microsimulation (PolicyEngine, 2025); here I apply it at the firm level. The generator takes the registration threshold as a parameter and is documented in Appendix A.1; I build one population for 2023–24 (threshold £85,000), the paper’s primary data year, and one for 2024–25 (threshold £90,000), used for the forward sweep of Section 4.

Construction. Index firms by i . For each sector s and ONS turnover band $b = [b, \bar{b})$, the *UK Business: Activity, Size and Location* table gives a business count $N_{s,b}$, and I draw $N_{s,b}$ firms with turnover inside the band. The within-band density is a maintained assumption, not an observed distribution. In the bottom band ($[0, 50)$ thousand pounds) the draw is uniform. In each closed band above it the draw is a truncated power law, $f(y) \propto y^{-\alpha_b}$ on $[b, \bar{b})$, with α_b the log–log slope of the band-average densities of the two neighbouring bands (a central difference, one-sided at the ends), so the density is monotone inside every band and approximately continuous across band edges; the exponents are 0.07, 1.21, 1.68, 1.79, 1.84 for the bands from £50,000 upward. The open “£5m and over” band is drawn uniform in log turnover on $[\$5m, \$50m)$. Band counts are preserved exactly under every choice; a uniform fill, used in earlier drafts, left a visible density step at each ONS band edge, including one at £100,000 inside the region the reform costings depend on. I interpret integer-labelled bands as contiguous half-open intervals: for example, the published “50–99” band is represented as $[50, 100)$ thousand pounds. The counts come from Table 8 of the ONS workbook (enterprises by SIC division and turnover sizeband, United Kingdom); the employment-band targets below come from Table 3 of the same workbook (enterprises by employment sizeband), so both ONS margins count the same statistical unit.

Two universes. The model distinguishes two populations and declares which one each calibration target describes. The *frame* is the ONS register of enterprises operating VAT and/or PAYE—about 2.72 million in 2024—which the population and employment targets describe. The *registered subset* is the population of HMRC VAT traders—about

2.18 million in 2023–24—which the HMRC counts by turnover band and by trade sector and the net-liability totals describe. The two differ in both directions. Above the threshold only about two-thirds of frame enterprises are VAT traders (from 54% in the £85,000–£150,000 band to 94% in the £1m–£10m band and 82% above £10m): the remainder are PAYE-only employers and businesses in exempt sectors (finance, health, education, and similar) outside the VAT net. The two top bands depend on how the open ONS “£5m and over” band is filled, since it has no published top; I draw it uniform in log turnover on [\$5m, \$50m), the maximum-entropy choice for an open band, and the band’s liability total is calibrated regardless, so reform costings near the threshold are unaffected. Below the threshold about 89% of frame enterprises are registered, largely because small businesses enter the register *through* VAT registration. HMRC also reports about 216,500 traders with negative or zero turnover who are not identifiable in the ONS turnover bands; I append them as an out-of-frame stratum before calibration. Neither universe is the full business population of about 5.5 million in the DBT business population estimates (Department for Business and Trade, 2024): the 2.9 million businesses registered for neither VAT nor PAYE are outside the frame. The difference is concentrated below the threshold and matters for threshold *cuts*, which I flag where relevant. Each firm is given input expenditure $x_i = \delta_i y_i$, with the input share δ_i drawn from a rescaled Beta distribution with sector-specific shifts, clamped to $\delta_i \in [0.1, 0.95]$ (mean about 0.6, so value added averages about 40% of turnover, in line with the UK non-financial business economy). Its net VAT liability is the standard rate applied to its value added,

$$v_i = \tau (y_i - x_i), \quad \tau = 0.20,$$

so every firm’s net rate $v_i/y_i = \tau(1 - \delta_i)$ lies in [1%, 18%], strictly positive and below the statutory cap of 20%. The model therefore abstracts from net-repayment traders (firms with zero-rated outputs whose input reclaim exceeds output VAT—HMRC’s negative-liability column); representing them requires rated-structure detail (zero, reduced, and exempt output shares) that the public aggregates do not resolve by band. Employment is assigned from the ONS employment-band shares conditional on each firm’s sector, and the same assignment is retained through calibration and final output. Appendix A.1 records the model and data provenance relevant to replication.

The base population reproduces the ONS structure but not the HMRC VAT-registered totals, so I re-weight it. Each firm receives a strictly positive weight $w_i = e^{\theta_i}$, and for a vector of official targets $T = (T_k)_{k=1}^K$ the weighted synthetic counterpart of target k is

$$\hat{T}_k(\theta) = \sum_i a_{ki} w_i,$$

where a_{ki} is firm i ’s contribution to target k on the universe the target describes. For the frame targets—the ONS enterprise total and the employment bands— a_{ki} is a membership indicator for frame firms and zero for the appended HMRC traders. For the registered-subset targets—HMRC trader counts by turnover band and by sector, and net liability by band— a_{ki} is the firm’s *registration propensity* $p_{b(i)}$ times the band or sector indicator (times v_i for a liability target), where p_b is the HMRC trader count in turnover band b

divided by the frame’s mass in that band; appended traders carry $p = 1$. At unit weights the registered-subset rows therefore reproduce the HMRC counts by construction and the frame rows reproduce the ONS counts by construction; the optimiser’s work is to fit the sector composition, the liability totals, and the near-threshold shape while holding both. I choose θ to minimise a symmetric relative-error loss with per-target importance weights λ_k ,

$$L(\theta) = \frac{1}{K} \sum_k \lambda_k \min\left\{(\hat{T}_k/T_k - 1)^2, (T_k/\hat{T}_k - 1)^2\right\} + \frac{\alpha}{N} \sum_i |\theta_i|,$$

the final term a mean absolute penalty on the log-weights. The targets are the HMRC VAT-registered counts by turnover band (all eight bands, including the negative-or-zero band) and by trade sector, the HMRC net VAT liability by turnover band, the ONS enterprise total and employment-band totals, and—for the 2023–24 vintage—the OBR’s published near-threshold profile. HMRC’s sector table includes traders whose turnover band is “Unknown” (159,300 in 2024–25, none in 2023–24); I allocate them proportionally so the sector and band rows count the same registered population. Chart C of the March 2023 *Economic and Fiscal Outlook* reports HMRC business counts in £1,000 turnover bands from £65,000 to £90,000 ([Office for Budget Responsibility, 2023](#)). I interpolate these counts to 2023–24 between the 2019–20 outturn and the OBR’s 2025–26 frozen-threshold projection. The final £90,000 row lies beyond the fine window and is not used.

All 25 retained bins enter as *shape* targets, applied separately on each side of the threshold: the below-threshold profile is normalised over £65,000–£85,000 and the above-threshold profile over £85,000–£90,000, and each is scaled to the frame’s own mass in its window. Chart C counts every business in HMRC records, a universe that below the threshold is wider than the ONS frame, so its levels are not imported on either side; the OBR data supply only the within-side geometry—the rise into the threshold and the decline beyond it—and the cross-threshold step comes from the frame’s own band structure. The resulting frame density steps from about 14,600 firms per £1,000 in the £84,000–£85,000 bin to 11,500 just above it (–21%), against the chart’s own –60%, because the chart’s below-threshold universe includes the unregistered businesses the frame does not. No fine bands are published for the £90,000 threshold era, so the 2024–25 vintage keeps coarse bands only. Turnover-band and near-threshold targets carry the largest λ_k , followed by VAT liability by band. Two liability dimensions are reported as informational diagnostics but excluded from the optimizer. VAT liability by sector is excluded because the generator draws a common input-share distribution with sector shifts rather than calibrating the input/output tax structure by sector. The below-threshold liability band (£1–threshold, £1.46bn in 2023–24) is excluded because that total is remitted by voluntary registrants—input-reclaim traders averaging about £2,150 of net remittance—whom the model’s standard-rate-on-value-added liability does not represent (Section 2); imposing it on the whole below-threshold population would force the optimiser to suppress weights in the £50,000–£85,000 region, distorting the near-threshold shape the fine targets exist to pin. The problem is solved by gradient descent, and the symmetric loss keeps targets of very different magnitudes on a common scale.

Scope and registration. After calibration each firm receives two flags. *Scope* marks firms in the VAT net: above the threshold, a seeded weighted selection within each HMRC turnover band whose cumulative weight reaches that band’s HMRC trader count (the remainder are out-of-scope PAYE-only or exempt-sector enterprises, which remit nothing under any schedule considered here); below the threshold, every frame enterprise is treated as registrable, a maintained assumption since HMRC publishes only the registered count there. *Registration* marks in-scope firms above the threshold and, below it, a seeded weighted selection reaching the HMRC £1-threshold count; appended negative-or-zero traders are in scope and registered. Registration above the threshold is thus certain for in-scope firms. This is an approximation to the statutory test, which is a rolling 12-month test on *taxable* turnover with a roughly two-month registration lag, a forward-looking 30-day test, and an exception for temporary breaches (Section 2); exempt-sector businesses are represented only through the out-of-scope share, not identified individually. For a single-year cross-section calibrated to registered-population totals these timing and composition wedges are second order for the aggregates, but they are worth naming: the HMRC band tables classify traders by total declared outputs (VAT-return Box 6), and the ONS turnover measure is likewise total turnover, so the running variable throughout is total rather than taxable turnover. The seeded selections match each weighted total to within one firm’s weight but do not identify *which* firms are in scope or register voluntarily; their composition remains a modelling assumption. The resulting dataset comprises approximately 2.94 million firm records per vintage (2.72 million frame draws plus 0.22 million appended traders), weighted to represent about 2.72 million frame enterprises in 2023–24, of which about 2.18 million are VAT-registered traders. Checked against the official sources along every calibrated dimension, the vintages attain overall calibration accuracies of 98.0% (2023–24) and 98.1% (2024–25); Appendix A.1 reports the full table. The headline is the mean of five calibrated-dimension scores, each clipped at zero as $\max\{0, 1 - |\hat{T} - T|/|T|\}$, HMRC dimensions scored on registered rows and ONS dimensions on frame rows; it excludes the VAT-liability-by-sector diagnostic described above. Because the turnover-band rows hold by construction, the informative residuals are the sector composition (93.3% and 92.6%) and the liability totals by band (98.4% and 98.9%).

Aggregate fit does not imply record-level representativeness. Kish effective sample sizes are 97,977 (3.3% of rows) and 225,945 (7.7%), while maximum weights are about 5,649 and 1,985; the median weight is one. The log-log within-band fill removed the extreme weights (above 2,000) that the uniform fill had needed in the large-turnover bands. Synthetic rows should therefore be read as weighted accounting support points, not independent firm observations; local and sub-band results can be sensitive to a small number of high-weight records. I report rounded scenario estimates and expose the full weight diagnostics in the generated calibration report.

The registration threshold in the data. My primary data year is 2023–24, when the statutory threshold was £85,000 (Figure 3 shows the calibrated distribution). The HMRC band targets step down across the threshold—678,350 firms in the £1–£85,000 band against 305,320 in £85,000–£150,000, or roughly 8,000 against 4,700 firms per £1,000 of band width—and the registered subset of the calibrated population inherits that band-

average step through the registration propensity; the frame density itself steps only from about 9,000 to 8,700 enterprises per £1,000 across the same bands. The *local* shape within a few thousand pounds of the threshold is pinned by the OBR near-threshold targets described above: the calibrated frame density rises into the threshold (about 14,600 firms per £1,000 in the £84,000–£85,000 bin), steps down across it (about 11,500 just above), and declines through the pinned window—the administratively observed bunching geometry (Office for Budget Responsibility, 2023; Liu et al., 2021), present in the file because it is targeted, not because any synthetic firm chooses it; no sub-band feature of band-calibrated synthetic data carries behavioural information beyond what its targets put there (Section 5 quantifies the resulting specification sensitivity). Counterfactual policies—including the £85,000 to £90,000 reform and the forward threshold sweep of Section 4—are evaluated on these firms by re-applying the registration rule to the microdata for the static costings, holding the empirical turnover distribution fixed. A behavioural layer that re-prices firms under counterfactual schedules, conditional on an assumed turnover elasticity, is developed in Section 7.




Figure 3: Synthetic firm turnover distribution (weighted firm counts, £1,000 bins), 2023–24. The near-threshold profile—rising into the £85,000 threshold and stepping down across it—is calibrated to the OBR’s published £1,000-band counts and is therefore target-inherited, not behavioural (Section 5). Levels are the ONS VAT/PAYE frame: the unregistered businesses that supply most of the cross-threshold cliff in all-business administrative charts are outside it, so level comparisons to such charts are not like-for-like. The steps at £65,000 and £90,000 (the edges of the OBR-pinned window) and at £150,000 (an HMRC band edge) are calibration-target boundaries, not features of the underlying firm distribution. The £90,000 boundary lies inside the bunching estimator’s fit range; Appendix A.2 measures its effect on the headline statistics.

4 Static costing of threshold reforms

A static costing measures the first-order revenue effect of a reform by re-applying the tax rule to a fixed firm population, with no behavioural response. Let firm i have turnover y_i , calibrated net VAT liability $v_i = \tau(y_i - x_i)$ (Section 3), weight w_i , scope flag s_i (one for firms in the VAT net), data-year registration flag q_i , and baseline voluntary-registration flag r_i (one for firms registered with data-year turnover at or below the data-year threshold T_0). A firm remits VAT only once registered. Under a counterfactual threshold T an in-scope firm is registered when its data-year turnover is at or above T ; a voluntary registrant ($y_i < T_0$) is unchanged, since no threshold move changes its obligation; out-of-scope firms never remit. Total VAT revenue at threshold T is therefore

$$R(T) = \sum_i w_i v_i s_i \mathbf{1}\{y_i \geq T \vee r_i\},$$

which reproduces the baseline registered set at $T = T_0$. Three conventions follow and are worth stating. First, because voluntary registrants already remit, a threshold cut credits only the liability of in-scope firms newly drawn into the net: about 89% of below-threshold frame enterprises are already registered (Section 3), and the unregistered businesses a cut would mainly reach lie outside the ONS frame (Section 8). Second, a threshold rise releases every in-scope registrant with data-year turnover in $[T_0, T)$: they no longer need to register, which is the reading HMRC’s costing uses; the statutory deregistration threshold (£2,000 below the registration threshold) would keep an existing registrant in the top of the band in the net, and I report that as a sensitivity. Third, the headline treats every released firm as deregistering; the anchor comparison reports the voluntary-retention sensitivity. Band membership is evaluated on *data-year* turnover in every fiscal year: a static cross-section has no growth process, and ageing the near-threshold stock across the threshold by a uniform factor would treat bunched firms as crossers, which they are not by construction (Section 5). Liabilities are aged by a cumulative nominal-growth factor—1.031, 1.052, 1.078, 1.110, and 1.142 for 2024–25 to 2028–29 relative to 2023–24 (the anchor runs on the 2023–24 vintage; the sweep on the 2024–25 vintage, aged relative to its own data year); the factors are assumptions. The estimate is the mechanical reclassification effect and excludes the behavioural response of Section 7. At the £90,000 threshold this puts the in-scope above-threshold VAT base at £183.4bn in 2025–26 and £188.7bn in 2026–27.¹

Anchor comparison against HMRC’s published costing. I compare the static pipeline against HMRC’s costing of the April 2024 increase from £85,000 to £90,000, an-

¹These aggregates are weighted sums of the in-scope registered population’s net VAT liabilities above the threshold, aged forward—a theoretical-liability (VTTL-style) base, not the published cash-receipts line (around £169bn in 2023–24; OBR forecast near £180bn for 2025–26), which is reduced by refunds, the VAT gap, and collection timing. Below-threshold voluntary remittances are excluded from the reported base because the model’s standard-rate-on-value-added liability is not calibrated for them (about £2.4bn model-implied against a £1.46bn HMRC total). The unaged bases are £175.0bn in 2023–24 and £180.5bn in 2024–25; the liability-by-band calibration scores (98.4% and 98.9%) summarise the remaining discrepancy against HMRC’s band totals. Band-to-band *differences*, which drive every costing in the paper, are less exposed than the level.

nounced and costed at Spring Budget 2024 (HM Treasury, 2024). The comparison requires the costing baseline convention: the official pre-measure baseline is not a frozen £85,000 threshold but a counterfactual path in which the freeze legislated to March 2026 ends and the threshold resumes uprating (85, 85, 87, 89, 92 thousand pounds over 2024–25 to 2028–29 in the model’s implementation), so the measure’s impact shrinks over the forecast and changes sign in 2028–29, when the counterfactual threshold overtakes £90,000. The official profile is $-150, -185, -125, -50, +65$ £m. Each year I difference revenue under the £90,000 policy and under that year’s counterfactual threshold, both evaluated with the registration rule above on data-year turnover and liabilities aged to the year. The mechanical headline, in which every released firm deregisters, gives $-198, -202, -123, -42, +88$: it reproduces the sign, the narrowing profile and the 2028–29 turning point, and lands within 9% of the official figure in 2025–26 and 1% in 2026–27, overshooting by 32% in 2024–25 (HMRC’s first year carries a registration lag the static rule has no counterpart for), undershooting by 15% in 2027–28, and overshooting the sign-flip year by 36% (where the result depends on the assumed £92,000 counterfactual threshold). No retention or behavioural parameter is tuned to produce this: the released population is the in-scope registrants with data-year turnover in the vacated band, about 28,600 firms in the first two years against HMRC’s expected 28,000 fewer registrants, and its liability is the calibrated standard rate on value added. Retaining the Liu et al. (2021) voluntary-registration share (43%) of released-firm liability gives $-113, -115, -70, -24, +88$, undershooting the official profile by 25–52%: LLAT’s share is a *firm* share among all below-threshold eligibles, and applying it as a liability share among just-released firms double-counts the voluntary registrants the rule already holds in the net, so it is a generous retention assumption rather than a central case. Ageing turnover by the same factors as liability, so that membership is evaluated on aged turnover, gives $-323, -376, -217, -66, +130$, between 1.3 and 2.2 times the official figures: it carries the near-threshold stock across the threshold, and the sign, narrowing profile and turning point—which follow from the counterfactual threshold path taken from the official costing convention—survive under either convention, so it is the level, not the shape, that the comparison tests. Keeping existing registrants above the £88,000 deregistration threshold in the net gives $-120, -122, -125, -43, +87$, a stock reading that does not follow the counterfactual path. Because my microdata are calibrated to HMRC VAT aggregates, even a close match confirms internal consistency rather than out-of-sample accuracy; HMRC’s own note expects about 28,000 fewer registrants in 2024–25 and 14,000 fewer on average over the forecast, against about 28,000 in-scope registrants in the released band here (Figure 4).

A menu of threshold locations. I next cost a menu of further threshold locations, taking the post-reform £90,000 threshold as the baseline. This sweep runs on the 2024–25 vintage population (Section 3). Each row of Table 1 costs a move of the threshold to a new location for 2025–26 by the registration rule above, holding the data-year turnover distribution fixed, so the revenue change reflects only the reclassification of in-scope firms into or out of the VAT net. Because the optimiser determines the within-band shape near the threshold (Section 3), and because the 2024–25 vintage has no fine-band targets, so that its density around £90,000 is the log-log within-band fill rather than a calibrated




Figure 4: Anchor reform: revenue impact of raising the VAT registration threshold from £85,000 to £90,000 (the April 2024 change), my estimate under the full-deregistration and 43% retention conventions versus HMRC’s published costing, 2024–25 to 2028–29. Both series narrow over the forecast and turn positive by 2028–29 because the costing baseline is a counterfactual threshold path that resumes uprating after March 2026 and overtakes £90,000. The mechanical estimate, with no retention or behavioural parameter, lies within a third of the official profile in each of the four release years (within 10% in 2025–26 and 2026–27) and 36% above it in the sign-flip year; the 43% retention convention undershoots it. Because the microdata are calibrated to HMRC and OBR aggregates, this is a consistency check, not an out-of-sample validation.

profile (Section 5), the $\pm \$5,000$ rows are conditional on that synthetic within-band allocation. Lowering the threshold draws in in-scope frame firms not already registered; raising it releases registered firms in the vacated band. The sweep is strongly asymmetric: a £20,000 cut to £70,000 raises only about £197m, because most frame firms just below the threshold are already registered, while a £30,000 rise to £120,000 costs about £1.27bn, 0.69% of the £183.4bn base. The cut rows are frame-confined lower bounds: the businesses a lower threshold would mainly reach are the unregistered sole traders outside the ONS frame, whom the data do not contain. Figure 5 shows the sweep graphically. These are *static* results: they bound the first-order fiscal effect of a threshold change but abstract from the behavioural response of Section 7.

Table 1: Static revenue and VAT-paying-firm effects of moving the VAT registration threshold, 2025–26, relative to the post-reform £90,000 baseline (in-scope above-threshold VAT base £183.4bn; 2024–25 vintage population). Figures are mechanical, holding data-year turnover fixed (no behavioural response); in-scope firms only. Voluntary registrants are held registered. Cut rows are frame-confined lower bounds (see text). This sweep is distinct from the anchor reform of Figure 4. The machine-readable table is `results/static_sweep.txt` in the replication repository.

Threshold	Revenue change (£m)	Change (%)	VAT-paying firms (000s)
£70,000	+196.8	+0.86	+246.3
£75,000	+151.6	+0.64	+175.3
£80,000	+103.6	+0.42	+110.9
£85,000	+53.7	+0.20	+52.7
£90,000 (baseline)	0	0	0
£95,000	−199.7	−0.11	−28.8
£100,000	−404.1	−0.22	−57.4
£105,000	−624.9	−0.34	−86.3
£110,000	−842.0	−0.46	−113.7
£115,000	−1,057.8	−0.58	−139.6
£120,000	−1,270.2	−0.69	−163.2

4.1 Costing schedule reforms

The threshold sweep above moves the *location* of a single cutoff. The same static principle—re-applying a tax rule to the fixed turnover distribution, with no behavioural response—also costs reforms that change the *shape* of the schedule rather than merely its position. I consider four reforms and report their first-order revenue effects against a common baseline: the £85,000-notch / £175.0bn 2023–24 base of in-scope firms. (This 2023–24 data-year base differs from the £183.4bn figure in the sweep above, which runs on the 2024–25 vintage aged to 2025–26 at the post-reform £90,000 threshold; the schedule reforms are costed in the data year to match the bunching and dominated-region analysis.




Figure 5: Static effect of moving the VAT registration threshold, 2025–26, relative to the £90,000 baseline (dashed line). Panel (a): revenue impact—cuts raise little within the frame because most below-threshold enterprises are already registered; rises release registered firms. Panel (b): the change in the number of VAT-paying firms, tracking the local density of firms on each side of the threshold.

Both bases are theoretical-liability (VTTL-style) aggregates in the sense of the footnote above, not cash-receipts lines.) For each reform I recompute the weighted sum of liabilities $R = \sum_i w_i v_i$ under the new schedule, holding each firm’s turnover fixed, and difference it against the baseline.

Threshold relocation. Raising the threshold from £85,000 to £100,000 releases the firms with turnover in $[\$85,000, \$100,000)$ —about 94,600 in-scope firms—from the VAT net, for a static revenue change of $-\text{£}590\text{m}$ on the common base (the menu applies each schedule to in-scope firms at their data-year turnover, without the deregistration-gap convention of the anchor). Because the released band lies *above* the data-year threshold, its liabilities are directly calibrated (and its lower reach pinned by the OBR fine-band targets), so the band is summed directly; a smooth-counterfactual alternative that fits the above-threshold per-£1,000 profile and integrates it across the band gives $-\text{£}615\text{m}$ and 90,500 firms—within 6% of the direct sum now that the within-band fill is monotone. On this common base the level move relocates the dominated region to the new cutoff and widens it to £25,000 (Section 6) at a static cost of $-\text{£}590\text{m}$; the linear taper removes the dominated region entirely at the menu’s largest static cost of $-\text{£}1,407\text{m}$, and a constant-marginal-rate taper of the same width does so at $-\text{£}932\text{m}$ (Section 6). Which matters more depends on what the reader wants from the reform—the dominated-region metric indexes misallocation only, and Section 8 discusses what it omits.

Graduated taper. A graduated taper replaces the discrete notch with a *marginal* remittance rate that phases in over a band, so liability is the integral of the marginal rate and a firm’s *effective* rate rises continuously from 0% to the full 20% across the band. The band top is a property of the chosen shape (Section 6): a taper whose marginal rate is constant at $m > \tau$ meets the standard-rate line at $U(m) = mT^*/(m - \tau)$, which is £141,667 at $m = 50\%$ and falls towards £106,250 as $m \rightarrow 100\%$, while a taper reaching

the full rate at £105,000 would push net revenue inside the band below its value at the threshold, recreating a dominated interval smoothly rather than at a cliff. I cost two designs on the same £85,000–£141,667 band: the *linear* taper, whose marginal rate rises from 0% to 100% across the band, and a *constant* 50% marginal-rate taper. Firms in the band remit the tapered schedule on their (fixed) turnover; mechanically the linear design reduces revenue by −£1,407m and the constant-rate design by −£932m, with about 278,100 in-scope firms in the band facing a lower effective rate. The linear design gives more relief low in the band at the price of a marginal remittance rate that reaches 100% at the top; the constant-rate design caps the marginal rate at 50% and costs a third less. A revenue-matched comparison makes the trade-off concrete: a constant marginal rate of 71.1% over £85,000–£118,200 removes the dominated region at exactly the static cost of the raise to £100,000 (£590m), reaching about 188,700 in-scope firms against the 94,600 the raise releases.

Banded reduced rate. A banded reduced rate charges firms in the £85,000–£105,000 band a reduced rate $\tau_{\text{low}} \in \{10\%, 15\%\}$ in place of the standard 20%, with 20% retained above £105,000. Statically this reduces revenue by −£398m at $\tau_{\text{low}} = 10\%$ and by −£199m at $\tau_{\text{low}} = 15\%$, with about 117,300 in-scope firms in the band facing the lower rate.

Table 2: Static, behaviour-free revenue effect of the reforms relative to the common £85,000-notch / £175.0bn 2023–24 baseline of in-scope firms. Each reform is costed mechanically by re-applying its schedule to the fixed turnover distribution, with no behavioural response.

Reform	Lever	Static (£m)
Raise threshold to £100,000	Location	−590
Linear taper, £85k–£141.7k	Shape (marginal rate 0 → 100%)	−1,415
Constant 50% marginal taper, £85k–£141.7k	Shape (marginal rate 50%)	−938
Reduced rate 10%, £85k–£105k	Rate (step)	−401
Reduced rate 15%, £85k–£105k	Rate (step)	−201

These static figures bound the first-order fiscal effect of each reform but abstract from the behavioural response of Section 7. Two maintained conventions carry through every costing: the firm bears the full statutory rate on its value added (no pass-through to prices), and per-firm liability is the standard rate on value added, so Flat Rate Scheme users—who are concentrated in exactly the £85,000–£150,000 bands—are represented only through the calibrated band totals, not individually. One feature of the shape reforms is not visible in a static costing: unlike a threshold move, which transports the notch to a new cutoff, the taper and the reduced rate *change the size of the notch itself*—the discrete drop in net revenue at the threshold and the empty band of turnover it induces. Quantifying that change requires the structural object I develop in Section 6, where I derive the notch’s dominated region and show analytically how each of these reforms alters it.

5 Bunching and the limits of band-calibrated synthetic data

This section applies the polynomial counterfactual-density estimator of [Chetty et al. \(2011\)](#) and [Kleven and Waseem \(2013\)](#), as used for UK VAT by [Liu et al. \(2021\)](#), to the synthetic population. Its finding is about the data, not about firms: every sub-band feature of a band-calibrated synthetic population is what its targets put there, so the excess mass the estimator reads at £85,000 is the OBR profile the population was calibrated to.

5.1 Methodology

Firms are binned in £1,000 intervals and a counterfactual density $g(y)$ is fitted by polynomial regression on bin counts outside a manipulation window around the threshold y^* , then projected into the window ([Appendix A.2](#) gives the specification). The upper edge of the window, y_R , is located by integrating missing mass upward from the threshold until it equals the excess mass below; where the missing mass inside the search window falls short, y_R is censored at the window edge and reported as such. Excess mass is the gross positive-part gap

$$E = \int_{y^*-W}^{y^*} [f(y) - g(y)]_+ dy,$$

reported alongside its signed net counterpart. The normalised excess b_{LLAT} divides the positive-part gap over $[y^* - W, y_R]$ by the average counterfactual height over that region, in the spirit of [Liu et al. \(2021\)](#) without reproducing their estimator exactly. No elasticity is computed: the identification critique of bunching designs applies ([Blomquist et al., 2021](#); [Bertanha et al., 2023](#)), and on these data there is no behavioural mass to convert.

5.2 Results: diagnostics, not behavioural estimates

At the £85,000 threshold ([Figure 6](#)) the baseline specification returns a gross excess mass of $E = 1,335$ weighted firms, a signed net excess of $-26,574$, a fitted missing mass of $\Delta_R = 2,476$ with an interior marginal buncher at $y_R = \$86,260$, and $b_{\text{LLAT}} = 0.178$ on the frame density. Nothing here is behavioural: the profile the estimator reads is the OBR £1,000-band profile the population was calibrated to, the generator contains no location choice, and the statistics are unstable across the specification grid (E from zero to about 8,300; b_{LLAT} from zero to 2.7). On in-scope firms alone the same specification gives $E = 8,668$ ([Appendix A.2](#)). The 2024–25 population, which has no fine-band targets, makes the complementary point: at its £90,000 threshold the estimator returns $E = 0$ and $b_{\text{LLAT}} = 0.000$: with no calibrated profile, the log–log within-band fill is smooth through the £90,000 HMRC band edge, the signed net excess is negative, and the estimator reads no false-positive bunching. The absence of any calibrated sub-band structure there is why the $\pm \$5,000$ rows of [Table 1](#) are flagged as conditional on the within-band allocation. [Appendix A.2](#) reports two constructed-data exercises that characterise the estimator itself.




Figure 6: Observed synthetic firm density and polynomial no-bunching counterfactual at the 2023–24 registration threshold. The 2023–24 population reproduces the administratively observed bunching profile because it is calibrated to the OBR £1,000-band targets; the estimator reads the target-inherited excess back as $E = 1,335$ ($b_{\text{LLAT}} = 0.178$). This is not behavioural: the generator has no firm location-choice, so all sub-band structure is inherited from the targets and the weight optimiser.

Comparison with the literature. Liu et al. (2021) (hereafter LLAT) document genuine bunching at the UK VAT notch in HMRC micro-data over 2004–05 to 2014–15, with an excess-bunching ratio of $b = 1.361$ (SE 0.202) and a manipulation window of $-\$14,000$ to $+\$24,000$. The synthetic $b_{\text{LLAT}} = 0.178$ is not a second estimate of their parameter: it is a readback of the OBR target profile. What does carry across is institutional context: LLAT report that about 43% of eligible below-threshold firms register voluntarily, which motivates the voluntary-retention sensitivity of Section 4.

6 The notch and its dominated region

The behavioural response to the registration notch is documented in administrative data by Liu et al. (2021); the synthetic data contain no bunching signal to exploit (Section 5). To characterise the notch precisely I therefore turn to its analytic structure, which yields an exact distortion measure that needs no behavioural calibration. I set out the firm-level structure following Kleven and Waseem (2013), Saez (2010), and Kleven (2016), whose bunching framework Liu et al. (2021) adapt to UK VAT. The UK registration threshold is a *notch*, not a kink: once turnover crosses the threshold the firm becomes liable for VAT on its *entire* base, not merely on the increment above it. The section’s central deliverable is the notch’s *dominated region*—an exact, elasticity-free property of the schedule that indexes the distortion the threshold imposes (Section 6.2)—and its analytic response to each schedule reform, complementing the static reform costs of Section 4.1. Section 7 prices the intensive-margin response to each reform conditional on an assumed turnover elasticity; nothing in this section depends on that assumption.

Scope. The model prices the registration margin for firms that remit positive net VAT. Two populations are outside its scope and are treated as such throughout: below-threshold *voluntary* registrants (roughly 43% of eligible firms in Liu et al., 2021), who have chosen registration and so face no notch at the threshold, and net-repayment traders (zero-rated exporters and similar), whose remittance is negative for rated-structure reasons the generator does not represent (Section 3).

6.1 The firm’s problem

Each firm i chooses turnover y_i to maximise after-tax profit. Let $c_i(y_i)$ denote the cost of producing turnover y_i , T^* the registration threshold, and τ the standard VAT rate. In the simplest statement VAT applies to the firm’s entire turnover once it registers, so after-tax revenue is discontinuous and profit is

$$\pi_i(y_i) = \begin{cases} y_i - c_i(y_i), & y_i < T^* \quad (\text{unregistered}), \\ (1 - \tau)y_i - c_i(y_i), & y_i \geq T^* \quad (\text{registered}). \end{cases} \quad (1)$$

At T^* after-tax revenue falls discretely by τT^* : a firm that nudges its turnover from just below to just above the threshold loses τT^* of net revenue on turnover it was already earning. With the standard rate $\tau = 0.20$ and threshold $T^* = \$85,000$ this discrete drop

is $\tau T^* = \$17,000$ on whole turnover. Real VAT taxes value added, and the corrected data give each firm net liability $v_i = \tau(y_i - x_i)$ with deductible inputs $x_i = \delta_i y_i$ (Section 3); profit is then $(1 - \delta_i)y_i - \tilde{c}_i(y_i)$ unregistered and $(1 - \delta_i)(1 - \tau)y_i - \tilde{c}_i(y_i)$ registered—the deductible share multiplies *both* branches, so the notch geometry of (1) carries over firm by firm, with the cash size of the jump scaled to $\tau(1 - \delta_i)T^*$ (about £6,800 at the mean value-added share of 40%). This drop—absent from a kink, where only the *marginal* rate changes at the threshold while average liability moves continuously—is the defining feature of a notch: it makes locating just below the cutoff strictly attractive and creates the empty band derived next. Throughout, the firm is assumed to bear the full statutory rate on its value added—no pass-through to prices—so the notch geometry is stated in terms of the firm’s own net revenue; with partial pass-through the cash size of the jump, and the width derived next, would shrink in proportion to the share borne by the firm.

6.2 The dominated region

The notch creates a range of turnover just above T^* that no profit-maximising firm chooses under the model. A firm just below the threshold earns net revenue T^* ; a registered firm at $T^* + a$ earns $(1 - \tau)(T^* + a)$. Before these net revenues become equal, the registered firm receives less net revenue while producing more. Thus, for any non-decreasing production cost, it earns strictly lower profit; no locally constant-cost approximation is required. Equating net revenues defines the boundary of this guaranteed dominated region:

$$(1 - \tau)(T^* + a) = T^* \implies a = T^* \frac{\tau}{1 - \tau}. \quad (2)$$

This is the Kleven–Waseem guaranteed dominated-region width (Kleven and Waseem, 2013). It holds the firm’s VAT-exclusive price fixed and lets the firm bear the whole τ ; if instead the consumer price is held fixed, so that the firm remits $\tau/(1 + \tau)$ of its receipts, the width is $\tau T^* = \$17,000$. The no-pass-through convention is used throughout and every width below scales by $(1 - \tau)$ under the other. At my parameters it gives $a = \$85,000 \times 0.20/0.80 = \$21,250$, so the dominated region is $(T^*, T^* + a) = (\$85,000, \$106,250)$, shown in Figure 7.

Two properties make this band a transparent index of the distortion. First, it is *exact and elasticity-free*, and—less obviously—*invariant to the firm’s input share*. Under the value-added formulation above, the indifference condition at the edge of the region is $(1 - \delta_i)(T^*) = (1 - \delta_i)(1 - \tau)(T^* + a)$: the deductible share $(1 - \delta_i)$ cancels, so every firm with positive value added faces the *same* dominated width $a = T^* \tau / (1 - \tau)$, whatever its input intensity. (A per-firm width built from the net remittance rate, such as $T^* \tau_{0i} / (1 - \tau_{0i})$, corresponds to a different technology in which marginal turnover carries no input cost, and is not used here.) The width depends on τ and T^* alone—not on the cost specification beyond monotonicity, any behavioural elasticity, or any smoothing parameter. Who escapes the notch is a matter of scope, not of input shares (the invariance is a property of the formulation: unregistered profit carries the same deductible share, so no input VAT goes unreclaimed below the threshold; if deregistered firms instead bore VAT on their inputs, the width would shrink in the input share and vanish for shares above one half—the

input-reclaim margin that drives voluntary registration, Section 2): voluntary registrants escape by choice, and net-repayment traders have no positive remittance to jump (the Scope paragraph above). Second, the region is an order of magnitude wider than the $\sim \$1k$ region implied by a smoothed *kink* approximation, and it spans a populous part of the firm distribution: about 129,900 weighted in-scope firms in the calibrated synthetic population lie in the $\$85,000$ – $\$106,250$ band (`results/dominated_region_mass.txt`; out-of-scope enterprises face no notch and are excluded). Because the generator has no location choice, this counts the firms the data *place* in the band—indexing the size of the population exposed to the distortion—rather than firms that optimally locate there; the count is conditional on the log-log within-band fill inside the coarse ONS bands and on the scope flag (Section 3) and would shift under a different within-band shape. It is the width a , not the count, that is exact.

The region under unreclaimed input VAT. The invariance above holds because the value-added formulation lets an unregistered firm deduct its inputs. In the UK system an unregistered firm cannot reclaim input VAT: it bears τ on its standard-rated purchases, which is the mechanism behind voluntary registration (Section 2). With input share δ and inputs fully standard-rated, unregistered net revenue is $y(1 - \delta - \tau\delta)$ against $(1 - \tau)(1 - \delta)y$ registered, and the indifference condition gives

$$a(\delta) = T^* \frac{\tau(1 - 2\delta)}{(1 - \tau)(1 - \delta)}, \quad a(\delta) = 0 \text{ for } \delta \geq \frac{1}{2}. \quad (3)$$

The width is firm-specific, falls in the input share, and vanishes once inputs exceed half of turnover, where registration pays for itself through reclaim. Two things follow. First, δ must be the *standard-rated* input share: labour is not an input carrying VAT and rent, financial and insurance purchases are exempt, so only a fraction s of the generator’s total input share δ_i bears reclaimable VAT. Second, $a(\delta) \leq 0$ is the model’s voluntary-registration mechanism: a firm for which registration pays at fixed turnover registers below the threshold. That gives an external test. With all inputs standard-rated ($s = 1$) the model predicts that 71% of below-threshold in-scope firms prefer registration, against the 43% that Liu et al. (2021) observe; the two reconcile at $s^* = 0.77$, a standard-rated share of purchases that is plausible for the consumer-service firms that dominate the near-threshold population but is not externally measured here. At s^* , 54% of in-scope firms within $\pounds15,000$ below to $\pounds45,000$ above the threshold face a positive dominated width (mean $\pounds8,129$ among them, $\pounds4,395$ across all of them), and about 25,000 weighted firms sit inside their own region, against the 129,900 in the fixed $\pounds21,250$ band (`results/dominated_region_mass.txt`). The $\pounds21,250$ figure is the zero-input, zero-pass-through upper bound. The reform comparisons change in one important way: the secondary notch at a reduced-rate band top is a step between two *registered* states, so its width a' is δ -invariant and every firm faces it. For the 46% of near-threshold firms with $a(\delta_s) = 0$, a banded reduced rate therefore *creates* a dominated interval at the band top where none existed, while a level rise changes nothing for them; the ordering stated below holds for the firms with a positive primary width. The behavioural layer of Section 7 adopts the same input-VAT formulation for the extensive margin: the marginal buncher falls in δ and does not exist above one half.




Figure 7: The VAT notch and the firm distribution. Panel (a) shows net revenue before production costs: immediately above £85,000, VAT reduces it from y to $0.8y$, making turnover in $(\$85,000, \$106,250)$ strictly dominated. Panel (b) shows the three-bin moving mean of weighted 2023–24 in-scope synthetic firm counts, with the same interval shaded. The distribution is descriptive and target-inherited, not the outcome of behavioural optimisation; its drop at £85,000 reflects the scope flag—every frame business below the threshold is registrable, while about half of frame enterprises above it are VAT traders—not an empty turnover interval. The vertical lines mark the marginal-buncher *ability* n_H at $e = 0.17$ for deductible shares $\delta = 0$ and $\delta = 0.6$ (Section 7); a firm of that ability registers at $n_H[(1 - \delta)(1 - \tau)]^e$, so the lines are not observed-turnover locations. The exact dominated width is $a = \$21,250$.

How each reform changes the dominated region. Because the width $a = T^*\tau/(1 - \tau)$ depends on the threshold and the rate alone, the effect of each schedule reform of Section 4.1 on the dominated region is exact and follows directly from the notch arithmetic of (2) and (4), with no dependence on the cost specification beyond monotonicity, any behavioural elasticity, or any simulated re-optimisation. Raising the threshold to £100,000 leaves the rate unchanged and therefore *relocates* the region—and, because the width is proportional to the threshold, *widens* it: at $T^* = \$100,000$, $a = \$100,000 \times 0.20/0.80 = \$25,000$, so the dominated region becomes $(\$100,000, \$125,000)$. Only the proportional width $a/T^* = 25\%$ is unchanged; the absolute band grows by £3,750. A banded reduced rate does *not* remove the distortion. It reduces the *largest single* notch—the jump at T^* falls from $\tau T^* = \$17,000$ to $r T^* = \$12,750$ at a 15% band rate—but, because the band reverts to the standard rate τ at its upper edge $T_1 = \$105,000$, it introduces a *secondary* notch there, and hence a second dominated region the simple “ a shrinks with τ ” reasoning omits. The primary region at T^* (the band rate r on whole turnover) has the usual width $a = T^*r/(1 - r)$: $\$85,000 \times 0.15/0.85 = \$15,000$ at 15% and $\$85,000 \times 0.10/0.90 = \$9,444$ at 10%. The secondary region at the band top—where the schedule steps from r back to τ on whole turnover—has width

$$a' = T_1 \frac{\tau - r}{1 - \tau}, \quad (4)$$

giving $\$105,000 \times 0.05/0.80 = \$6,562$ at 15% and $\$105,000 \times 0.10/0.80 = \$13,125$ at 10%. The *total* dominated-turnover width is therefore $a + a' = \$21,563$ at 15% and $\$22,569$ at 10%—both marginally *above* the $\$21,250$ baseline, not below it. The reduced rate splits one wide empty band into two narrower ones rather than shrinking the total: it lowers the largest single notch (to $\$12,750$ at 15%) at the price of adding a secondary notch of $(\tau - r)T_1 = \$5,250$. The firm counts confirm this (`results/dominated_region_mass.txt`): the primary regions hold about 94,600 (15%) and 60,300 (10%) weighted in-scope firms, the secondary regions add about 34,000 and 65,000, and the totals—128,600 and 125,300—lie within 3.5% of the 129,900-firm baseline band. The total dominated *width* is larger under both reduced rates and the exposed population is essentially unchanged: the reduced rate relocates the distortion, it does not remove it. A graduated taper, by contrast, can remove the dominated region entirely—provided its schedule keeps net revenue monotone, and that condition is restrictive. Any band-confined schedule whose effective rate reaches the full τ by £105,000 leaves net revenue $(1 - \tau) \times \$105,000 = \$84,000$ at the band top, below the £85,000 available at the threshold itself, so net revenue must fall over some range inside the band—a smoothly dominated interval in place of the cliff. Monotonicity requires the marginal remittance rate $m(y)$ to stay below one, and band confinement requires liability to meet the standard-rate line τy at the band top U . For a constant marginal rate $m > \tau$ these two conditions pin $U(m) = mT^*/(m - \tau)$: the band top is a property of the schedule’s *shape*, falling from £141,667 at $m = 50\%$ towards the infimum $T^*/(1 - \tau) = \$106,250$ as $m \rightarrow 100\%$, where net revenue is flat across the band and the region is only weakly dominated. The linear taper costed in Section 4.1 phases m from 0 to 100% over $[\$85,000, T^*/(1 - 2\tau)] = [\$85,000, \$141,667]$; the constant-rate taper costed alongside it holds $m = 50\%$ over the same band. Both leave liability continuous at

both band edges and net revenue non-decreasing throughout, so the band of dominated turnover, with the discrete incentive to suppress turnover, is removed entirely. The price of removal is a wide band and, for the linear design, a marginal remittance rate that reaches 100% at its top; the constant-rate design caps the marginal rate at 50% and costs a third less. On the dominated-region metric, then, only the taper removes the distortion; the reduced rate splits it and the level move relocates and widens it. The metric indexes the misallocation distortion only—what it omits, and how the reforms compare on cost, is taken up in Section 8. The dominated region is the paper’s exact, elasticity-free measure of the distortion the notch imposes, and the analytic statement above is exactly how much of that distortion each reform of Section 4.1 removes.

7 Behavioural reform costs

The costings of Section 4 hold turnover fixed. This section adds an intensive-margin response for an *assumed* turnover elasticity e , reported as a sensitivity range that nests the static costing exactly at $e \rightarrow 0$. Its result is deflationary: on the intensive margin the static numbers are robust.

7.1 An iso-elastic response, region-confined

Following Kleven and Waseem (2013) and Saez (2010), a firm of ability n_i chooses turnover y to maximise

$$\pi_i(y) = (1 - \delta_i)(1 - \tau f(y))y - C(y; n_i, e), \quad C(y; n, e) = \frac{n}{1 + 1/e} \left(\frac{y}{n}\right)^{1+1/e}, \quad (5)$$

where δ_i is the input share, $f(y) \in [0, 1]$ the fraction of the statutory rate the schedule levies at y , and C the own-factor cost (bought-in inputs are already netted out through $1 - \delta_i$). Under a flat rate the optimum is

$$y^* = n [(1 - \delta)(1 - \tau)]^e, \quad (6)$$

so e is the elasticity of turnover with respect to the retained wedge. Each firm re-optimises *within the region of the reform schedule that contains its observed turnover*, taking that turnover as optimal under the baseline notch; inverting (6) at the baseline fraction f_0 and re-solving at the reform fraction f_1 , ability and input share cancel:

$$y_i^* = y_{\text{obs},i} \left(\frac{1 - \tau f_1}{1 - \tau f_0} \right)^e, \quad (7)$$

clipped at the region’s edges. Relocation across a notch is the extensive margin, which the synthetic data cannot identify (Section 5) and which Section 6 characterises analytically; it is out of scope here, and confining the response to the firm’s own region is what makes the $e \rightarrow 0$ limit collapse onto the static costing exactly. The ratio is the same under the input-VAT formulation of Section 6, because a firm that stays in its region reclaims (or

bears) input VAT on both sides of the comparison; only the marginal buncher $n_H(e, \delta)$ differs between the two, and under the input-VAT formulation it falls in δ (£127,400, £124,900 and £116,800 at $e = 0.17$ for $\delta = 0, 0.2, 0.4$) and does not exist for $\delta \geq \frac{1}{2}$ (`results/dynamic_reform_results.txt`). The elasticity is not identified from these data; the sweep $e \in \{0.05, 0.17, 0.32\}$ spans the range used in the threshold literature, with 0.05 the Kleven and Waseem (2013) anchor. Liu et al. (2021) deliberately do not convert their bunching estimate into a turnover elasticity, so no published UK VAT-base elasticity anchors the sweep.

7.2 Results

Table 3 reports the behavioural cost of the three flat-rate reforms at each e alongside the static figure. Two features follow directly from the region-confined solve. A level rise has *no* intensive-margin offset at any e : released firms expand by $(1 - \tau)^{-e}$ (about +3.9% at $e = 0.17$) but have left the VAT base, so the loss is their full baseline remittance; firms above the new threshold face $f = 1$ before and after; and the only candidates for a taxed response—released firms whose expansion would carry them across the new notch—are clipped at it by construction. The invariance is therefore a property of the maintained status convention, not a computed statement about global optima. A reduced-rate band scales band firms up by (7) with $f_1 = r/\tau$ (about +2.0% at the 10% rate and $e = 0.17$), an offset of about £8m against a static £398m, reaching £14m at $e = 0.32$; Figure 8 traces the costs across e . The released code asserts on every run that the baseline solve reproduces observed turnover, that the $e \rightarrow 0$ cost equals the static cost to £0.1m, and that the solved response satisfies $d \ln y^* / d \ln(1 - \tau f) = e$.

Table 3: Behavioural revenue effect of the three flat-rate reforms relative to the £85,000-notch baseline, by assumed turnover elasticity e , with the static ($e \rightarrow 0$ limit) figure for comparison. Each behavioural figure re-solves every firm’s optimum within its schedule region per (7), holding e fixed. The final column reports the number of firms whose optimal turnover moves by more than £500 at $e = 0.17$. Estimates are conditional on the assumed e (Section 7.1). The graduated taper is excluded; see the paragraph below.

Reform	Static	$e = 0.05$	$e = 0.17$	$e = 0.32$	Firms ($e=0.17$)
Raise threshold to £100,000	−\$590m	−\$590m	−\$590m	−\$590m	91,796
Reduced rate 10% [\$85k, \$105k]	−\$401m	−\$399m	−\$393m	−\$387m	120,282
Reduced rate 15% [\$85k, \$105k]	−\$201m	−\$199m	−\$194m	−\$190m	120,282

Note. The tapers are omitted: their rate varies continuously with turnover, so their intensive response is outside the region-confined solve, and their marginal-rate channel is not represented by the flat-rate first-order condition. Their e -free figures—the static costs of −£1,407m (linear) and −£932m (constant 50%) and the exact removal of the dominated region—are in Sections 4 and 6. Firms are in-scope VAT firms with turnover in £70,000–£130,000, the range the region-confined solve re-optimises.

The intensive-margin layer is a robustness statement: naming an elasticity barely moves the static numbers, not at all for a level move and by under 6% for the reduced-




Figure 8: Behavioural reform cost against the assumed turnover elasticity e , for the three flat-rate reforms. Each curve is the change in revenue relative to the £85,000-notch baseline; the intercept at $e = 0$ is the static cost of Table 2. The raise-to-£100,000 line is exactly flat (no intensive-margin offset); the reduced-rate lines decline by under 6% of their static values across the plotted range. Dotted lines mark the swept values $e \in \{0.05, 0.17, 0.32\}$. Estimates are conditional on the assumed e (Section 7.1).

Table 4: Summary of headline estimates, consolidating results defined in the body sections cross-referenced below. Rows 1–2 are 2025–26 figures (the anchor on the 2023–24 population, the sweep on the 2024–25 population at its £90,000 baseline); the remaining rows refer to the £85,000 data-year baseline.

Object	Definition	Value	Source
Anchor, £85k to £90k, 2025–26	mechanical / 43% retention	–\$202m / –\$115m (HMRC –\$185m)	§4
Sweep, £90k baseline	£70k to £120k	+\$1.59bn to –\$1.28bn	Table 1
Dominated region	$a = T^*\tau/(1 - \tau)$; $a(\delta)$ with input VAT	£21,250 (B2C bound); 46% of firms	§6.2
In-scope firms in dominated band	weighted	$\approx 130,000$	§6.2
Raise to £100k	static = behavioural	–\$590m	§4, §7
Taper, £85k–£141.7k	linear / constant 50%	–\$1,407m / –\$932m	§4.1
Reduced rate, £85k–£105k	10% / 15%	–\$398m / –\$199m	§4.1
Bunching readback, 2023–24	E ; b_{LLAT}	1,335; 0.178 (target-inherited)	§5

rate bands. The behavioural uncertainty that could move these costings materially sits on the extensive margin—registration, deregistration and location choice at the notch—which this paper prices only through the analytic geometry of Section 6.

8 Conclusion

This paper contributes an open, auditable firm-level accounting microsimulation for analysing the UK VAT registration threshold, released in full—including its synthetic-data generator—for replication. Its value is as transparent policy-scenario infrastructure rather than a structural identification result: I make no claim to recover policy-invariant firm primitives, unlike Garicano et al. (2016) and Gourio and Roys (2014). The analysis is anchored to the 2023–24 data year, in which the statutory threshold was £85,000, with a second 2024–25 population (threshold £90,000) for the forward sweep. Table 4 collects the headline objects.

Findings. On the common £85,000 baseline, raising the threshold to £100,000 costs £590m, a taper across the £85,000–£141,667 band costs £1,415m (linear phase-in) or £938m (constant 50% marginal rate), and the 10% and 15% reduced-rate bands cost £401m and £201m. The mechanical anchor is consistent with HMRC’s published costing profile to within a third in each of the four release years (within 10% in 2025–26 and 2026–27) and 36% in the sign-flip year, once band membership is evaluated on data-year turnover; it is an internal consistency check rather than external validation, and the registration behaviour of released firms remains the decisive convention.

The notch creates an exact £21,250 dominated region for a firm that bears no unreclaimed input VAT; with input VAT borne, the region is firm-specific, zero for 46% of near-threshold firms and about £8,100 on average for the rest, so £21,250 is the business-to-consumer upper bound. Raising the threshold relocates and widens it, and a reduced-rate band divides it between two notches; only the continuous taper removes it. Intensive-margin responses change the reduced-rate costings only modestly under the assumed elasticities, while a threshold rise has no intensive-margin revenue offset in this model.

The bunching exercise supplies a methodological warning, not a behavioural estimate. Its 2023–24 profile is inherited from OBR calibration targets, its degree–window estimates are unstable. The administrative evidence in [Liu et al. \(2021\)](#) therefore remains the relevant evidence on firm behaviour.

Limitations and future work. The model omits location choice, voluntary-registration choice, compliance costs, sectoral incidence, pass-through ([Benedek et al., 2015](#); [Copestake et al., 2026](#)), and general-equilibrium effects. Its population is the ONS VAT/PAYE frame with an HMRC-calibrated registered subset: threshold cuts are costed only on frame enterprises, most of which are already registered, and omit the unregistered sole traders outside the frame that a lower threshold would mainly reach; extending the frame to the DBT business population is the natural next build. A fixed cross-section has no registration response or entry margin, which is why the anchor’s first year, where HMRC’s costing carries a registration lag, is the year it fits least well. The nominal-growth factors used for ageing are assumptions recorded in the replication repository. It also cannot capture persistent growth effects extending beyond the cutoff, as documented for a different payroll-tax notch by [Jysmä et al. \(2026\)](#). Those estimates do not transfer directly to UK VAT. Administrative firm-level validation, an explicit registration model, and welfare accounting are the main priorities for future work.

References

- Asatryan, Zareh and Andreas Peichl (2017) “Responses of Firms to Tax, Administrative and Accounting Rules: Evidence from Armenia,” CESifo Working Paper 6754, CESifo, https://ideas.repec.org/p/ces/ceswps/_6754.html.
- Bashir, Muhammad, Zehra Farooq, Usama Jamal, and Mazhar Waseem (2024) “Size-Based Policies and Firm Growth: Evidence from Pakistan,” working paper, University of Manchester, https://mazharwaseem.com/static/uploads/PAK_VAT_Threshold_October_2024.pdf.
- Benedek, Dora, Ruud A. de Mooij, Michael Keen, and Philippe Wingender (2015) “Estimating VAT Pass Through,” IMF Working Paper WP/15/214, International Monetary Fund, <https://www.elibrary.imf.org/view/journals/001/2015/214/001.2015.issue-214-en.xml>.
- Bertanha, Marinho, Andrew H. McCallum, and Nathan Seegert (2023) “Better Bunching, Nicer Notching,” *Journal of Econometrics*, 237 (2), 105512, <https://doi.org/10.1016/j.jeconom.2023.105512>.
- Best, Michael Carlos, Anne Brockmeyer, Henrik J. Kleven, Johannes Spinnewijn, and Mazhar Waseem (2015) “Production versus Revenue Efficiency with Limited Tax Capacity: Theory and Evidence from Pakistan,” *Journal of Political Economy*, 123 (6), 1311–1355, <https://doi.org/10.1086/683849>.
- Blomquist, Sören, Whitney K. Newey, Anil Kumar, and Che-Yuan Liang (2021) “On Bunching and Identification of the Taxable Income Elasticity,” *Journal of Political Economy*, 129 (8), 2320–2343, <https://doi.org/10.1086/714446>.
- Chetty, Raj, John N. Friedman, Tore Olsen, and Luigi Pistaferri (2011) “Adjustment Costs, Firm Responses, and Micro vs. Macro Labor Supply Elasticities: Evidence from Danish Tax Records,” *Quarterly Journal of Economics*, 126 (2), 749–804, <https://doi.org/10.1093/qje/qjr013>.
- Copestake, Alexander, Matthieu Bellon, and Wenzhang Zhang (2026) “Supply and Demand Determinants of Heterogeneous VAT Pass-Through,” [10.2139/ssrn.4207882](https://ssrn.com/abstract/4207882), National Tax Journal; publication metadata forthcoming.
- Council of the European Union (2020) “Council Directive (EU) 2020/285 on the Special Scheme for Small Enterprises,” Official Journal of the European Union, <https://eur-lex.europa.eu/eli/dir/2020/285/oj>.
- Crawford, Ian, Michael Keen, and Stephen Smith (2010) “Value Added Tax and Excises,” in Stuart Adam, Timothy Besley, Richard Blundell et al. eds. *Dimensions of Tax Design: The Mirrlees Review*, 275–362, Oxford: Oxford University Press.
- Department for Business and Trade (2024) “Business Population Estimates for the UK and Regions 2024: Detailed Tables,” National Statistics, October 2024, <https://www.gov.uk/government/statistics/business-population-estimates-2024>.
- Dharmapala, Dhammika, Joel Slemrod, and John Douglas Wilson (2011) “Tax Policy and the Missing Middle: Optimal Tax Remittance with Firm-Level Administrative Costs,” *Journal of Public Economics*, 95 (9–10), 1036–1047.

- Fossen, Frank M. and Stefan Bach (2008) “Reforming the German Local Business Tax—Lessons from an International Comparison and a Microsimulation Model,” *FinanzArchiv / Public Finance Analysis*, 64 (2), 245–272.
- Garicano, Luis, Claire Lelarge, and John Van Reenen (2016) “Firm Size Distortions and the Productivity Distribution: Evidence from France,” *American Economic Review*, 106 (11), 3439–3479, <https://doi.org/10.1257/aer.20130232>.
- Gourio, François and Nicolas Roys (2014) “Size-Dependent Regulations, Firm Size Distribution, and Reallocation,” *Quantitative Economics*, 5 (2), 377–416, <https://doi.org/10.3982/QE338>.
- Harju, Jarkko, Tuomas Matikka, and Timo Rauhanen (2019) “Compliance Costs vs. Tax Incentives: Why Do Entrepreneurs Respond to Size-Based Regulations?” *Journal of Public Economics*, 173, 139–164, <https://doi.org/10.1016/j.jpubeco.2019.02.003>.
- HM Revenue and Customs (2024) “Increasing the VAT Registration Threshold,” Tax information and impact note, GOV.UK, <https://www.gov.uk/government/publications/vat-increasing-the-registration-and-deregistration-thresholds/increasing-the-vat-registration-threshold>.
- HM Treasury (2018) “VAT Registration Threshold: Call for Evidence,” HM Treasury, March 2018, <https://www.gov.uk/government/consultations/vat-registration-threshold-call-for-evidence>.
- (2024) “Spring Budget 2024: Policy Costings,” HM Treasury, March 2024, https://assets.publishing.service.gov.uk/media/65e7920c08eef600155a5617/Published_Costing_Document_Spring_Budget_2024_Final.pdf.
- Jysmä, Sami, Youssef Benzarti, and Jarkko Harju (2026) “Policy Thresholds as Growth Barriers: Theory and Evidence from a Payroll Tax Notch,” NBER Working Paper35484, National Bureau of Economic Research, 10.3386/w35484.
- Kanbur, Ravi and Michael Keen (2014) “Thresholds, Informality, and Partitions of Compliance,” *International Tax and Public Finance*, 21 (4), 536–559.
- Keen, Michael and Jack Mintz (2004) “The Optimal Threshold for a Value-Added Tax,” *Journal of Public Economics*, 88 (3–4), 559–576, [https://doi.org/10.1016/S0047-2727\(02\)00165-2](https://doi.org/10.1016/S0047-2727(02)00165-2).
- Kleven, Henrik J. (2016) “Bunching,” *Annual Review of Economics*, 8, 435–464, <https://doi.org/10.1146/annurev-economics-080315-015234>.
- Kleven, Henrik J. and Mazhar Waseem (2013) “Using Notches to Uncover Optimization Frictions and Structural Elasticities: Theory and Evidence from Pakistan,” *Quarterly Journal of Economics*, 128 (2), 669–723, <https://doi.org/10.1093/qje/qjt004>.
- Liu, Li, Ben Lockwood, Miguel Almunia, and Eddy H. F. Tam (2021) “VAT Notches, Voluntary Registration, and Bunching: Theory and UK Evidence,” *Review of Economics and Statistics*, 103 (1), 151–164, <https://direct.mit.edu/rest/article/103/1/151/97767>.

- Liu, Li, Ben Lockwood, and Eddy H. F. Tam (2024) “Small Firm Growth and the VAT Threshold: Evidence for the UK,” IMF Working Paper WP/24/33, International Monetary Fund, <https://www.imf.org/en/publications/wp/issues/2024/02/16/small-firm-growth-and-the-vat-threshold-evidence-for-the-uk-544462>.
- Nandi, Tushar and Ross Warwick (2020) “How Businesses Respond to VAT: Evidence from West Bengal, India,” TaxDev / Institute for Fiscal Studies, <https://www.taxdev.org/research-projects/how-businesses-respond-vat>.
- Office for Budget Responsibility (2023) “The Impact of the Frozen VAT Registration Threshold,” Economic and Fiscal Outlook, March 2023, <https://obr.uk/box/the-impact-of-the-frozen-vat-registration-threshold/>.
- Office of Tax Simplification (2017) “Value Added Tax: Routes to Simplification,” OTS report, November 2017, https://assets.publishing.service.gov.uk/media/5a822577e5274a2e87dc1620/Value_added_tax_routes_to_simplification_web.pdf.
- Onji, Kazuki (2009) “The Response of Firms to Eligibility Thresholds: Evidence from the Japanese Value-Added Tax,” *Journal of Public Economics*, 93 (5–6), 766–775, <https://doi.org/10.1016/j.jpubeco.2008.10.003>.
- PolicyEngine (2025) “PolicyEngine UK: Open-Source Tax-Benefit Microsimulation,” <https://policyengine.org>.
- Saez, Emmanuel (2010) “Do Taxpayers Bunch at Kink Points?” *American Economic Journal: Economic Policy*, 2 (3), 180–212, <https://doi.org/10.1257/pol.2.3.180>.
- Seely, Antony (2024) “VAT: The Registration Threshold,” House of Commons Library Briefing Paper SN00963, <https://commonslibrary.parliament.uk/research-briefings/sn00963/>.

A Appendix

This appendix collects supporting material in two parts: data construction and calibration robustness (Appendix A.1); and statistical inference for the reduced-form bunching results (Appendix A.2), including the specification-sensitivity grid establishing that band-calibrated synthetic data cannot support bunching inference in either direction.

A.1 Data construction detail

Section 3 describes the two-stage synthetic-population pipeline, parameterised by the registration threshold, and its calibration to HMRC and ONS aggregates. This appendix records the implementation detail deferred from that account.

Each of the approximately 2.94 million firm rows carries a turnover, an input expenditure, a net VAT liability ($v_i = 0.20 (y_i - x_i)$), a calibration weight, a frame indicator, a VAT-scope flag, and a registration flag. The weights are chosen so the weighted frame rows reproduce the ONS enterprise total and employment bands while the weighted registered rows reproduce the HMRC trader counts and liability totals (Table A.1). Because the population is calibrated *to* the HMRC aggregates, agreement between my weighted outputs and those aggregates is a check of *internal consistency*, not external validation; the turnover-band rows in particular hold by construction through the registration propensity. The headline calibration accuracies (98.3% and 98.2% for the two vintages) and the anchor comparison of Section 4 are therefore consistency checks on the pipeline, not out-of-sample tests of firm behaviour.

Table A.1: Calibration accuracy by dimension, $\max\{0, 1 - |\hat{T} - T|/|T|\}$ averaged over the dimension’s targets. HMRC dimensions are scored on VAT-registered rows, ONS dimensions on frame rows. Source: `results/calibration_accuracy.txt`.

Dimension (universe)	2023–24	2024–25
HMRC turnover bands (registered)	100.0%	100.0%
ONS enterprise total (frame)	99.8%	99.9%
ONS employment bands (frame)	99.9%	99.9%
HMRC sector counts (registered)	93.4%	92.4%
HMRC net liability by band (registered)	97.1%	98.2%
Overall (five calibrated dimensions)	98.3%	98.2%
<i>Diagnostic:</i> net liability by sector (registered)	42.2%	44.0%
<i>Diagnostic:</i> net liability below threshold	31.0%	0.0%

The reported calibration score is a bounded summary statistic. For each dimension I compute $\max\{0, 1 - |\hat{T} - T|/|T|\}$, so a dimension missed by more than 100 percent receives zero rather than a negative score. The headline “overall” number is the unweighted mean of the five dimensions used as calibration targets—HMRC turnover bands, ONS population, employment bands, HMRC sector counts, and VAT liability by turnover band. VAT

liability by sector is excluded from that mean and reported separately as an informational diagnostic, because the current generator does not calibrate sector-specific input/output VAT structure.

The replication report also publishes the calibration-weight minimum, median, upper quantiles, maximum, coefficient of variation, and Kish effective sample size. These diagnostics matter because aggregate target agreement can coexist with a small number of influential synthetic records. The report is generated rather than transcribed, and the test suite verifies source-band support, sector-conditional employment assignment, and weighted voluntary registration.

Counterfactual exclusion window and polynomial degree. The no-VAT counterfactual density of Section 5 is fitted by polynomial regression on the observed density outside a manipulation window around the threshold. The baseline excludes $[T^* - 15k, T^* + 15k]$ and fits a degree-7 polynomial. Widening the window guards against contamination of the counterfactual by the behavioural response but reduces the data available for the fit; raising the degree improves in-sample fit but risks over-fitting the tails. The sensitivity of the excess-mass and bunching-ratio estimates to window width and degree, with associated uncertainty, is reported in Appendix A.2.

A.2 Inference for the bunching estimator

This appendix documents the inference procedure for the reduced-form bunching results of Section 5: the mass-conservation constraint that disciplines the counterfactual density, why no standard errors are reported, sensitivity to the polynomial degree and exclusion window. All numbers are produced by the released estimator run on the regenerated weighted synthetic 2023–24 population; the retained machine-readable artifact is `results/bunching_inference.txt`. The estimator reports the geometry only—the bunching ratio b , its Liu et al. (2021)-normalised counterpart b_{LLAT} , the excess mass E , and the mass-conservation objects (Δ_R, y_R) —and computes no elasticities: the mapping from excess mass to an elasticity is not identified without functional-form assumptions (Blomquist et al., 2021; Bertanha et al., 2023), and on these data there is no behavioural mass to convert in the first place.²

Mass-conservation constraint

The canonical bunching estimator of Chetty et al. (2011) and Kleven and Waseem (2013) imposes that the counterfactual density conserves mass: every firm removed from above the threshold reappears as excess mass below it. An unconstrained polynomial fit on bins outside a hand-set window computes the excess mass E (below T^*) and the missing mass Δ_R (above T^*) *independently*, with no guarantee that $E = \Delta_R$ and no endogenous location of the marginal buncher. I impose the constraint in two ways. First, the counterfactual

²Earlier versions of the estimator also printed a CES substitution elasticity and two wedge-normalised turnover elasticities, the latter inheriting an ad hoc $\tau/2$ normalisation from a deleted smoothed-schedule model. These outputs were labelled non-identified but were still reported; they have been removed from the code and the paper.

Table A.2: Headline bunching statistics (point estimates; degree-7 polynomial, $\pm\$15\text{k}$ window). The specification grid characterises sensitivity; standard errors are not reported—see the text.

Vintage	E	Δ_R	b_{LLAT}	y_R (£k)
2023–24 (£85k, OBR fine targets)	1,335	2,476	0.178	86.3
2024–25 (£90k, coarse bands only)	0	216	0.000	105.0 (censored)

polynomial in $(y - T^*)$ is fit by an iterated procedure in which the density of the bins above the excluded region is rescaled until the fitted counterfactual integrates to the observed total mass over the estimation range $[\$20\text{k}, \$140\text{k}]$ (the regressor is normalised to $[-1, 1]$ for numerical stability of the high-degree basis). Second, the upper edge of the manipulation region—the marginal buncher y_R —is located endogenously by integrating missing mass upward from T^* until it equals the excess mass below,

$$\int_{T^*}^{y_R} [f^{\text{cf}}(y) - f^{\text{obs}}(y)]_+ dy = \int_{T^*-W}^{T^*} [f^{\text{obs}}(y) - f^{\text{cf}}(y)]_+ dy = E.$$

On the 2023–24 population the baseline gives $E = 1,335$, $\Delta_R = 2,476$, and $y_R = \$86,260$, interior to the search window, so mass conservation binds. These mass-conservation outputs are estimator diagnostics, not evidence of a marginal behavioural buncher.

Why no standard errors are reported

Earlier drafts attached bootstrap standard errors to the bunching statistics (resampling firm rows with replacement and re-running the estimator). Those numbers are not reported anywhere in this paper, because the exercise has no estimand: the synthetic file is a deterministic construction from published targets, not a sample from a population, and row-resampling dispersion scales as $1/\sqrt{n}$ in a row count the analyst chooses. A standard error that can be made arbitrarily small by generating more synthetic rows conveys precision that does not exist. The specification grids below instead show how point estimates move with the estimator’s own choices. Generator Monte Carlo error is quantified separately, by rebuilding the full 2023–24 population under alternative generator seeds (`results/seed_sensitivity.txt`): across seeds the half-ranges are ± 156 on E , ± 0.01 on b_{LLAT} , $\pm \$4.2\text{m}$ on the raise costing, $\pm \$5.8\text{m}$ on the taper costing, and $\pm \$0.07\text{bn}$ on the in-scope base—the last larger than in earlier builds because the scope flag is itself a seeded draw within each turnover band. Table A.2 reports the point estimates. Two conventions matter for reading them. Bins are centred on integer £k values (edges at half-integers), so the bin containing the threshold spans £84.5k–£85.5k and its mass is attributed to the above side, attenuating E and Δ_R relative to edge-aligned binning. And E sums the *positive part* of the observed-minus-counterfactual gap in the excluded window below the threshold, so the statistic reads gross rather than net excess.

Sensitivity to polynomial degree and exclusion window

Table entries for a grid of polynomial degree $\in \{5, 6, 7, 8\}$ and symmetric exclusion window $\in \{\$10k, \dots, \$25k\}$ are in `results/bunching_inference.txt`; the pattern matters more than the cells. In the 2023–24 grid, E ranges from zero to about 190,500 firms and b_{LLAT} from zero to about 20.7. This specification sensitivity, together with the absence of a location-choice mechanism in the generator, precludes interpreting the baseline as a behavioural estimate.

Placebo and recovery exercises

Two constructed-data exercises characterise the estimator itself. On a placebo population whose density is reweighted to a smooth profile across £85,000, the estimator returns $E = 0$: it produces no false positive on a step-free world. When a Kleven–Waseem-consistent relocation of known size is injected into that placebo, the estimator recovers 0%, 12% and 31% of injected excess masses of 2,000, 5,000 and 8,000 firms, with mass conservation then binding at an interior y_R : the counterfactual is fitted on bins that include the depressed region beyond the exclusion window, so the estimator attenuates a genuine response rather than inflating it (`results/placebo_bunching.txt`, `results/recovery_bunching.txt`).